

The Dancing Screen

Independent Study
Mary Karroqe


Overview

- Project Description
- Original Approach
- Final Approach
- Data + Pre-Processing
- RNN
- Results
- Future Work

Project Description

The Problem

- My computer doesn't **dance**
- I wanted to **have fun** with my senior design project
- Not **saving the world** here



"Monster Dancing Sticker" via
[Giphy](#)

The Project

*Can I teach my computer to
generate dance videos?*

Applications

- Quarantine Dance Parties
- Looking cool in front of the children in your life
- Joy



"DISCO FEVER" Fortnite Dance

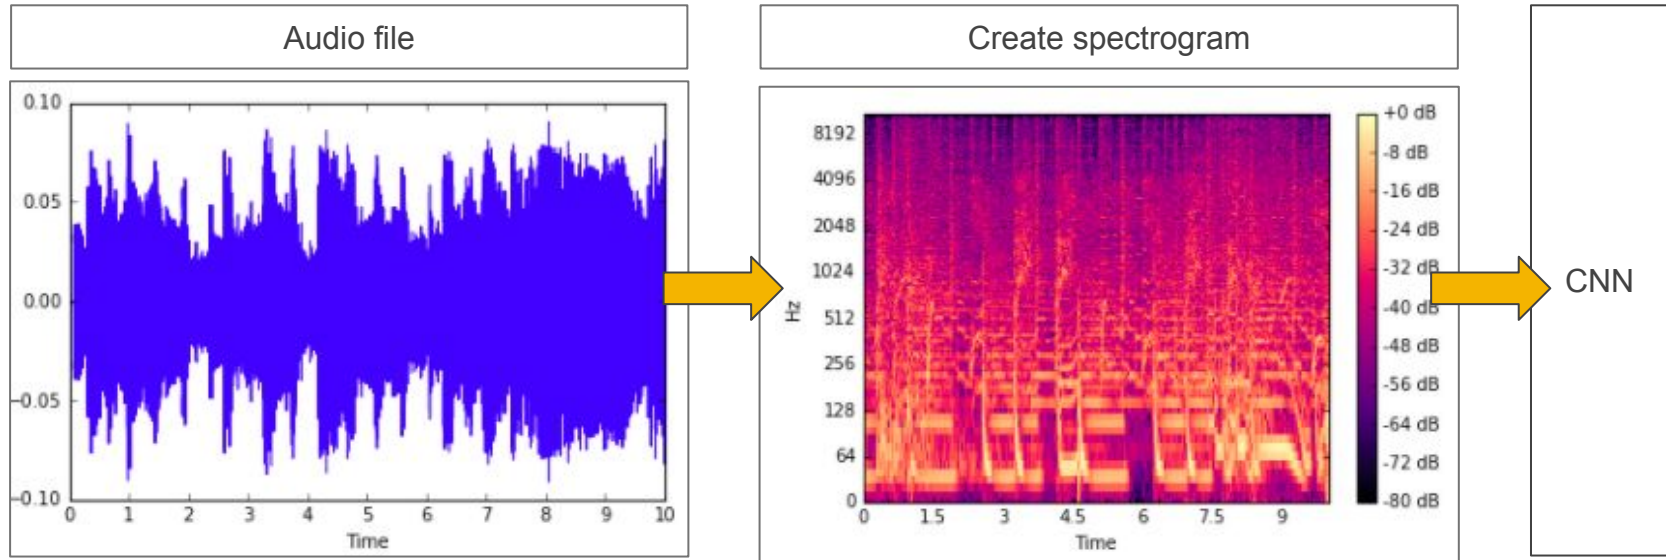
Applications

- Observing how the model breaks down dance moves
- Potential for dance instructors?
- Kids toys?




Original Approach

Original Approach



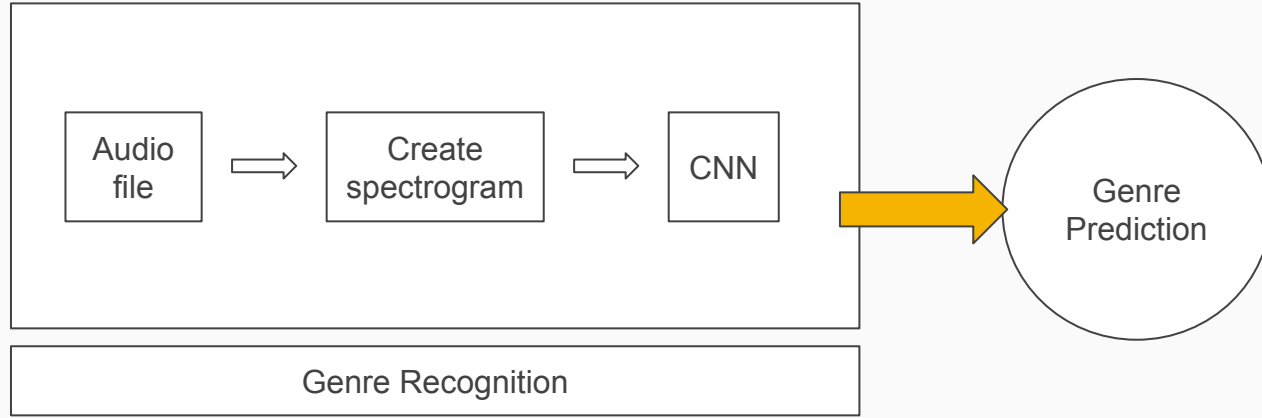
Genre Recognition

Original Approach

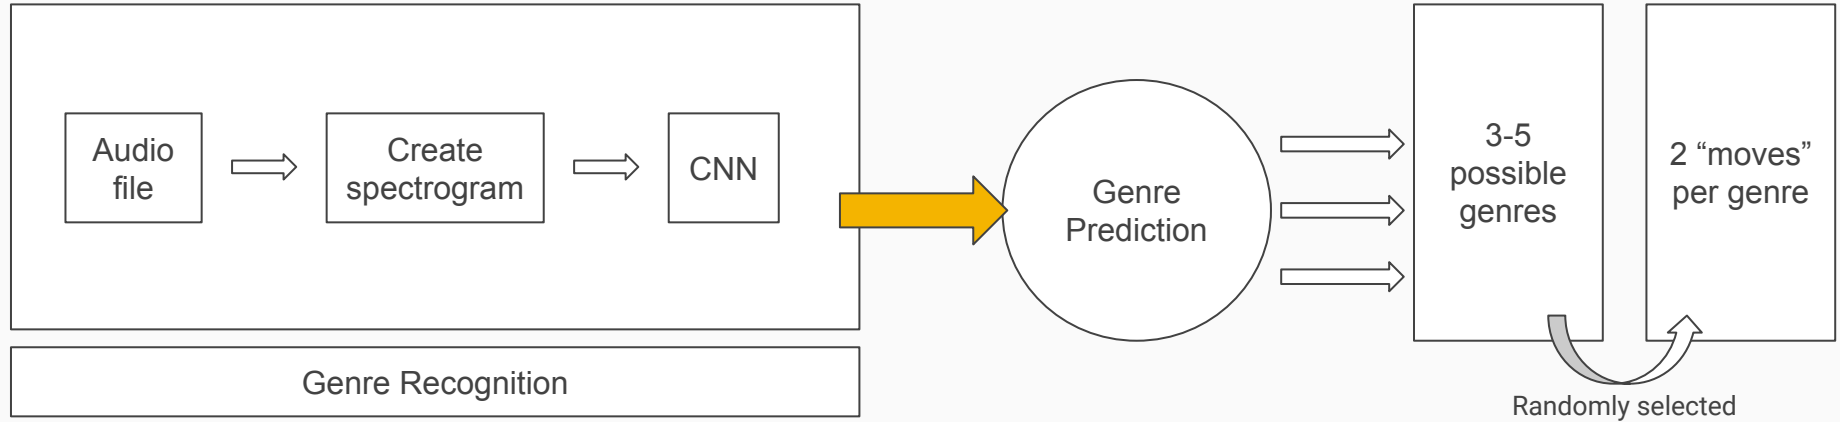


Genre Recognition

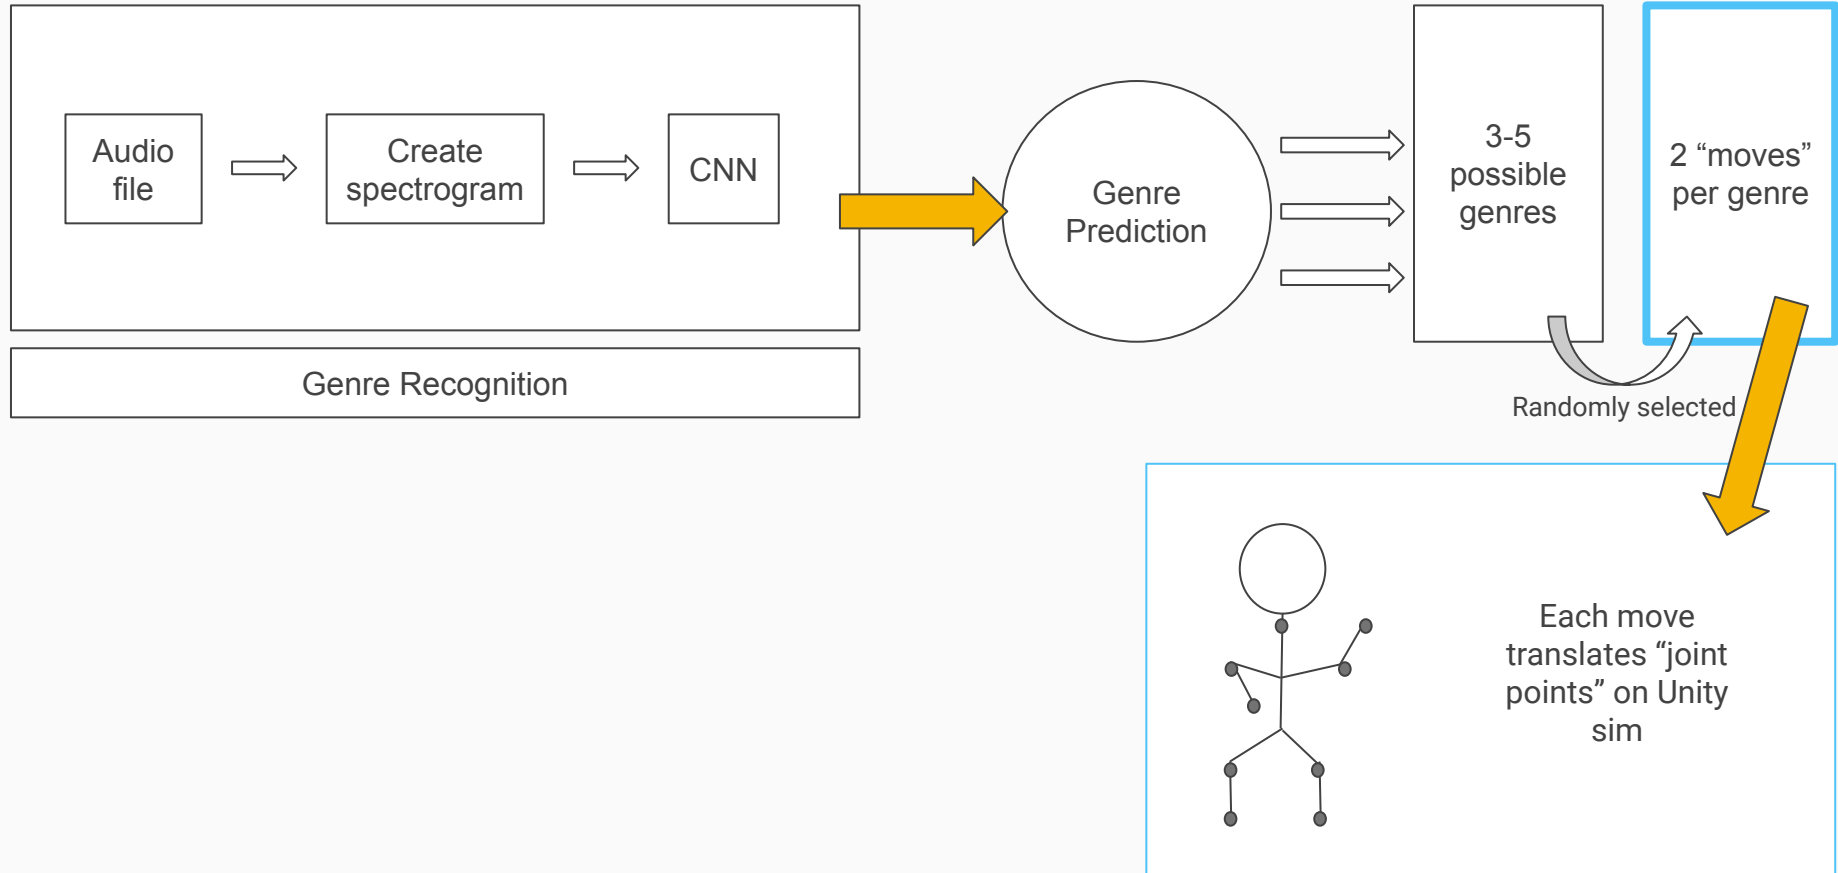
Original Approach



Original Approach



Original Approach




Problems with this Approach

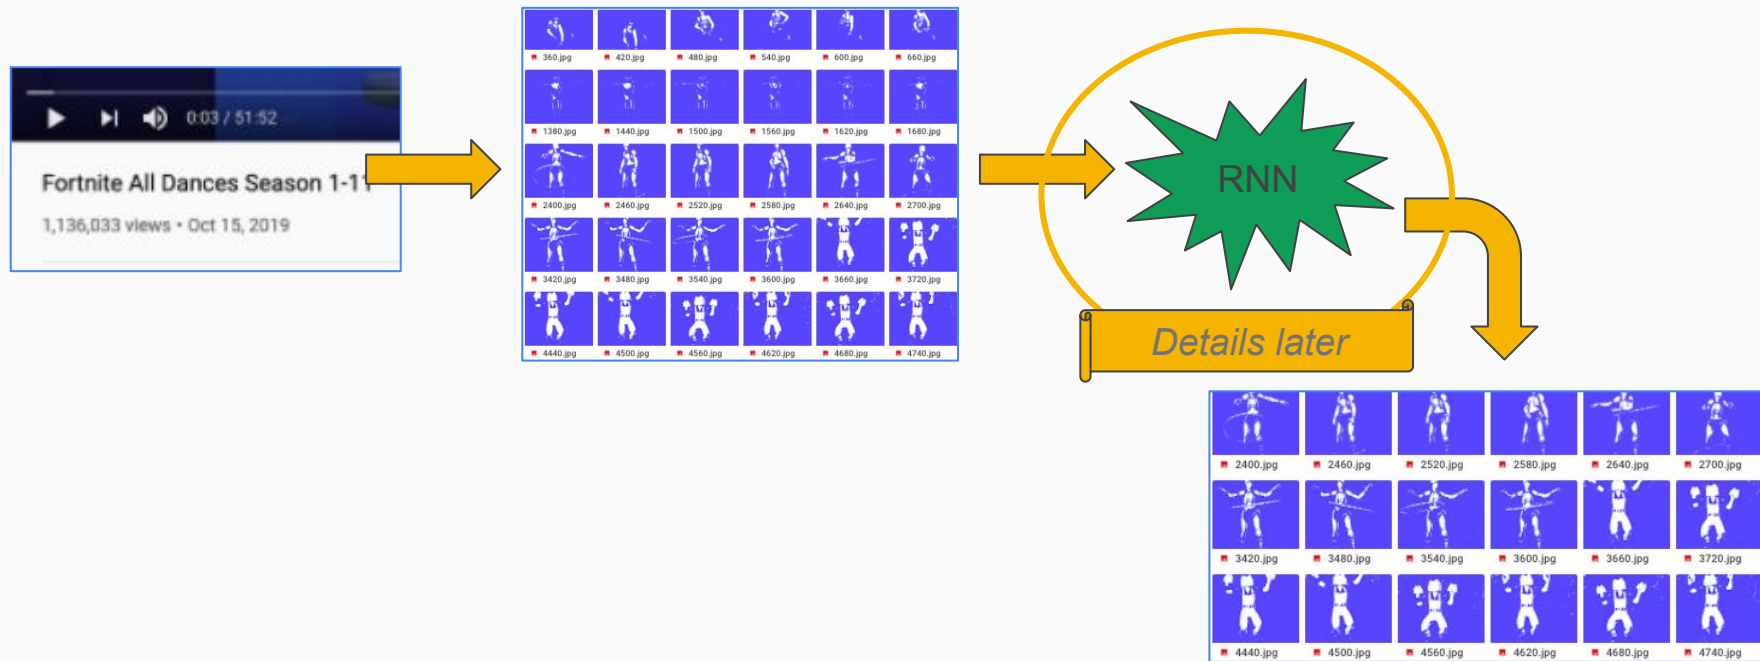
- The decision to start with genre recognition was **arbitrarily chosen**
- My genre recognition predictions from last semester were only **slightly more accurate** than random
- Convoluted pipeline; many aspects would be **hard-coded**
- **Limiting the output** would result in boring output
- Learning **Unity + connecting** to CNN output

Final Approach

Final Approach



Final Approach




Data + Pre-Processing

Input Data Options



Green Screen Silhouettes, YouTube playlist




"Shadow Dancers", YouTube video



Tik Tok videos filtered by hashtags of dance challenges

Data

- 51:52 Compilation Video
- 146 different dances
- Extracting 1 frame every 2 seconds
- 93,360 total images



Fortnite All Dances Season 1-11

1,136,033 views • Oct 15, 2019

7K 369 SHARE SAVE ...



deathmule
1.23M subscribers

SUBSCRIBE




DROOP

Dig deep.



DROOP

Dig deep.






DROOP

Dig deep.




DROOP


Dig deep.



Transformation 1: Increase contrast



Transformation 2: Remove blue pixels





Transformation 4 + 5: Convert all pixels to 1 color + crop






Sample dance move

Autoencoder

What are Autoencoders?

- Data compression algorithm
- Learned automatically from examples, usually with neural networks
- Data-specific
- Lossy
- Useful for dimensionality reduction



From [Keras Blog](#)

Encoder

- Used to *compress* original **images** (1,032 x 950) into (128) **vectors**, losing as little information as possible
- Create **dense representation** to be fed into RNN

Decoder

- Used to *expand* original (128) **vectors** into (1,032 x 950) **images**, *reconstructing* as much information as possible
- “**Translate**” RNN’s output

Version 1: Simple

- Single fully-connected layer
- As both encoder and decoder
- Image converted to greyscale
 - Single color channel

(703, 950, 1032)


(50, 950, 1032)

(703, 980400)
(50, 980400)







```
Epoch 45/50  
60000/60000 [=====] - 1s 14us/step - loss: 0.1061 - val_loss: 0.1042  
Epoch 46/50  
60000/60000 [=====] - 1s 14us/step - loss: 0.1057 - val_loss: 0.1038  
Epoch 47/50  
60000/60000 [=====] - 1s 13us/step - loss: 0.1052 - val_loss: 0.1034  
Epoch 48/50  
60000/60000 [=====] - 1s 13us/step - loss: 0.1048 - val_loss: 0.1030  
Epoch 49/50  
60000/60000 [=====] - 1s 13us/step - loss: 0.1045 - val_loss: 0.1026  
Epoch 50/50  
60000/60000 [=====] - 1s 13us/step - loss: 0.1041 - val_loss: 0.1023
```



10 epochs



50 epochs



150 epochs



They're the same picture.

Something's gotta change

```
Epoch 138/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 139/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 140/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 141/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 142/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 143/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 144/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 145/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 146/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 147/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 148/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 149/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
Epoch 150/150
703/703 [=====] - 3s 4ms/step - loss: 0.6931 - val_loss: 0.6931
<keras.callbacks.callbacks.History at 0x7f74b7daec50>
```

Something's gotta change

Mod 1:

- Realized my compression factor was off (~3%)
- Fixed it to be 24.5%

ResourceExhaustedError:

- It was able to handle `encoding_dim = 512`
- 19.14% compression
- Compiled, but while training:

```
ResourceExhaustedError: OOM when
      [[node gradients_4/loss_5
Hint: If you want to see a list of
      [Op:__inference_keras_scratch_gra

Function call stack:
keras_scratch_graph
```

Mod 2:


- Resize images before feeding
converting into np arrays
- 5% of original


- `(703, 47, 51)`

- Previously:


```
(703, 950, 1032)
```

```
(703, 980400)  
(50, 980400)
```





50 epochs



250 epochs

50, 250 epochs w scaled images




50 epochs

250 epochs

They're the same picture.



Mod 3:

- Time spent training clearly *not* the issue
- Changing the optimizer from rms_prop to `adadelta`
- 50 epochs



Mod 2 and Mod 3 Comparison

- Mod 2 more noisy
- Mod 3 has lightness in center where character should be




Version 2: Sparsity Constraint

- In V1, reps constrained only by size of hidden layer
- Learns approx of PCA
- Adding sparsity constraint


Layer (type)	Output Shape	Param #
input_6 (InputLayer)	(None, 2397)	0
dense_13 (Dense)	(None, 64)	153472
dense_14 (Dense)	(None, 2397)	155805

=====
Total params: 309,277
Trainable params: 309,277
Non-trainable params: 0

```
% add a dense layer with a L1 activity regularizer
encoded = Dense(encoding_dim, activation='relu',
                activity_regularizer=regularizers.l1(10e-5))(input_img)
```



50 epochs




100 epochs

Version 3: Deep Encoder

- Hypothesis:
more layers == more learning?

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 2397)	0
dense_1 (Dense)	(None, 512)	1227776
dense_2 (Dense)	(None, 128)	65664
dense_3 (Dense)	(None, 64)	8256
dense_4 (Dense)	(None, 32)	2080
dense_5 (Dense)	(None, 64)	2112
dense_6 (Dense)	(None, 128)	8320
dense_7 (Dense)	(None, 512)	66048
dense_8 (Dense)	(None, 2397)	1229661
Total params: 2,609,917		
Trainable params: 2,609,917		
Non-trainable params: 0		



100 epochs

Version 4: Convolutional

- Convolutional neural networks are pretty much standard when it comes to training on images, so it's application here makes sense

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	(None, 47, 51, 1)	0
conv2d_7 (Conv2D)	(None, 47, 51, 16)	160
conv2d_8 (Conv2D)	(None, 47, 51, 8)	1160
conv2d_9 (Conv2D)	(None, 47, 51, 8)	584
max_pooling2d_1 (MaxPooling2D)	(None, 24, 26, 8)	0
conv2d_10 (Conv2D)	(None, 24, 26, 8)	584
conv2d_11 (Conv2D)	(None, 24, 26, 8)	584
conv2d_12 (Conv2D)	(None, 22, 24, 16)	1168
conv2d_13 (Conv2D)	(None, 22, 24, 1)	145
Total params: 4,385		
Trainable params: 4,385		
Non-trainable params: 0		


```
-----  
ValueError
```

```
<ipython-input-68-40
```

```
3
```

```
4
```

```
----> 5
```



```
/usr/local/lib/python
```

```
143
```

```
144
```

```
--> 145
```

```
146
```

```
147
```

```
return d
```

```
ValueError: Error wh
```

```
SEARCH STACK OVERFLOW
```

Autoencoder Results

- Not a massive difference between different versions
- Wondering if there's

RNN

Why RNN?

- Feedforward networks have no memory
- To process a sequence, the network has to see it all as once
 - Turning sequence into single data point
- A dance is a sequence
 - Of moves, of frames
- A recurrent neural network processes sequences by iterating through the sequence elements and maintaining a state containing information relative to what it has seen so far

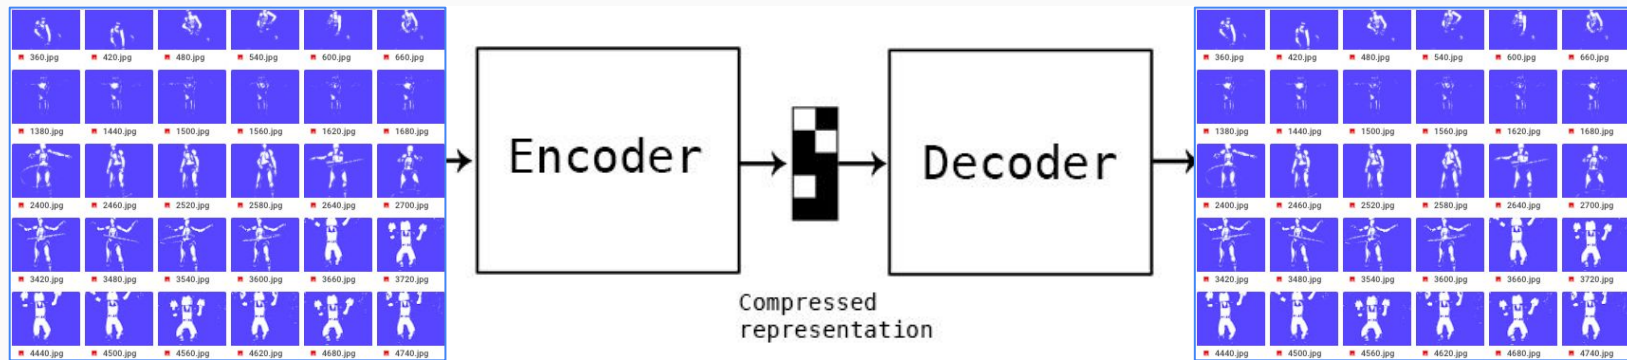
RNN

- Modifying Andrej Karpathy's RNN
 - Shakespeare creator
 - One-hot encoded characters
 - Input: 95-D vector representing 95 possible chars
 - *How can I do this with images? → binary images (two colors)*
 - *How can I pass in a*



RNN




- Compressing them normally is not good enough
 - Ex: regular JPEG compression algorithm is general
- We can use an autoencoder


Future Work




Final Approach









Our Baseline:



Future Work

- Incorporating audio files into input
- Preprocessing work has been completed

Thank you! 🙌

The Dancing Screen

Independent Study
Mary Karroqe