

.mp3 as input..

As I was playing with my first Colab notebook, I realized there was a problem with how I was setting up my model. The inputs I was prepared to feed in were the 518 attributes per track in the FMA dataset. In addition to being a very complex and large input, it's not how I want the model to function.

My goal is to input an .mp3 file and have the model output a genre prediction. I feel I have not framed my model/problem this way and will start fresh setting it up here.

Information from Audio Files

In the FMA dataset, the 9 features are the following:

Number of Instances	106 , 574 tracks	
Audio Features [9 total]	Feature Definition	Attributes per Feature [518 total]
Chroma	A quality of pitch class which refers to the "color" of a musical pitch (src)	84
Tonnetz	"Tonal grid;" planar array of pitches along three simplicial axes, corresponding to the three consonant interval (src)	42
Mel Frequency Cepstral Coefficient (MFCC)	<u>Cepstrum</u> : The info of rate of spectral bands <u>Mel scale</u> : scale that relates perceived frequency of a tone to the actual measured frequency MFCC features represent distinct units of sound as the shape of the vocal tract (src)	140
Spectral centroid	Characterizes spectrum (<u>where center of mass of the spectrum is located</u>). Connected to impression of brightness of sound (src)	7
Spectral bandwidth	Determines the resolution of a signal.. (src)	7
Spectral contrast	Level difference between peaks and valleys in the spectrum (src)	49
Spectral rolloff	Measure of the amount of the right-skewedness of the power spectrum (src)	7
Root Mean Square Energy	...have not found audio-related definition yet, just AC voltage/current here	7
Zero-crossing rate	The rate of sign-changes along a signal (src)	7

Extracting these features from an .mp3 using Librosa

Python's Librosa library includes many feature extraction methods, including the following spectral features:

<code>chroma_stft</code> ([y, sr, S, norm, n_fft, ...])	Compute a chromagram from a waveform or power spectrum.
<code>chroma_cqt</code> ([y, sr, C, hop_length, fmin, ...])	Constant-Q chromagram
<code>chroma_cens</code> ([y, sr, C, hop_length, fmin, ...])	Computes the chroma variant "Chroma Energy Normalized".
<code>melspectrogram</code> ([y, sr, S, n_fft, ...])	Compute a mel-scaled spectrogram.
<code>mfcc</code> ([y, sr, S, n_mfcc, dct_type, norm, lifter])	Mel-frequency cepstral coefficients (MFCCs)
<code>rms</code> ([y, S, frame_length, hop_length, ...])	Compute root-mean-square (RMS) value for each frame, either as a time series or as a single value.
<code>spectral_centroid</code> ([y, sr, S, n_fft, ...])	Compute the spectral centroid.
<code>spectral_bandwidth</code> ([y, sr, S, n_fft, ...])	Compute p'th-order spectral bandwidth.
<code>spectral_contrast</code> ([y, sr, S, n_fft, ...])	Compute spectral contrast [R6ffcc01153df-1]
<code>spectral_flatness</code> ([y, S, n_fft, hop_length, ...])	Compute spectral flatness
<code>spectral_rolloff</code> ([y, sr, S, n_fft, ...])	Compute roll-off frequency.
<code>poly_features</code> ([y, sr, S, n_fft, hop_length, ...])	Get coefficients of fitting an nth-order polynomial to the coefficients of the spectrogram.
<code>tonnetz</code> ([y, sr, chroma])	Computes the tonal centroid features (tonnetz), following the method of [1] .
<code>zero_crossing_rate</code> (y[, frame_length, ...])	Compute the zero-crossing rate of an audio time series.

These features correspond ***perfectly*** with what's in the FMA dataset; they're all here:

- chroma
- tonnetz
- mfcc
 - it also computes melspectrograms, which I hope can be a cool/useful visualization
- spectral_centroid
- spectral_bandwidth
- spectral_contrast
- spectral_rolloff
- rms (root mean square energy)
- zero_crossing_rate

Let's explore this data:

Helpful article link here (<https://towardsdatascience.com/extract-features-of-music-75a3f9bc265d>)

In [2]:

```
import librosa
audio_path = "audio_files/Don't Run With Pizzas.mp3"
x, sr = librosa.load(audio_path) # decodes audio file to 1D array of time series x, and sample rate of x, sr
print(type(x), type(sr))
```

```
/Users/mkarroqe/anaconda3/lib/python3.5/importlib/_bootstrap_external.py:415: ImportWarning: Not importing directory /Users/mkarroqe/anaconda3/lib/python3.5/site-packages/virtualenvwrapper: missing __init__
```

```
  _warnings.warn(msg.format(portions[0]), ImportWarning)
/Users/mkarroqe/anaconda3/lib/python3.5/importlib/_bootstrap_external.py:415: ImportWarning: Not importing directory /Users/mkarroqe/anaconda3/lib/python3.5/site-packages/sphinxcontrib: missing __init__
```

```
  _warnings.warn(msg.format(portions[0]), ImportWarning)
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/numba/errors.py:137: UserWarning: Insufficiently recent colorama version found. Numba requires colorama >= 0.3.9
```

```
  warnings.warn(msg)
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/librosa/core/audio.py:146: UserWarning: PySoundFile failed. Trying audioread instead.
```

```
  warnings.warn('PySoundFile failed. Trying audioread instead.')
```

```
<class 'numpy.ndarray'> <class 'int'>
```

In [3]:

```
# Playing audio  
import IPython.display as ipd  
ipd.Audio(audio_path)
```

```
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/IPython/core/  
formatters.py:92: DeprecationWarning: DisplayFormatter._ipython_dis  
play_formatter_default is deprecated: use @default decorator instea  
d.
```

```
def _ipython_display_formatter_default(self):  
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/IPython/core/  
formatters.py:98: DeprecationWarning: DisplayFormatter._formatters_  
default is deprecated: use @default decorator instead.
```

```
def _formatters_default(self):  
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/IPython/core/  
formatters.py:677: DeprecationWarning: PlainTextFormatter._deferred  
_printers_default is deprecated: use @default decorator instead.
```

```
def _deferred_printers_default(self):  
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/IPython/core/  
formatters.py:669: DeprecationWarning: PlainTextFormatter._singleto  
n_printers_default is deprecated: use @default decorator instead.
```

```
def _singleton_printers_default(self):  
/Users/mkarroqe/anaconda3/lib/python3.5/site-packages/IPython/core/  
formatters.py:672: DeprecationWarning: PlainTextFormatter._type_pri  
nters_default is deprecated: use @default decorator instead.
```

```
def _type_printers_default(self):
```

Out[3]:

0:00 / 3:38

In [4]:

```
#display waveform
%matplotlib inline
import matplotlib.pyplot as plt
import librosa.display
plt.figure(figsize=(14, 5))
librosa.display.waveplot(x, sr=sr)
```

```
/Users/mkarroge/anaconda3/lib/python3.5/site-packages/ipykernel/pyl
ab/config.py:66: DeprecationWarning: metadata {'config': True} was
set from the constructor. Metadata should be set using the .tag()
method, e.g., Int().tag(key1='value1', key2='value2')
```

```
    inline backend."""
```

```
/Users/mkarroge/anaconda3/lib/python3.5/site-packages/ipykernel/pyl
ab/config.py:44: DeprecationWarning: InlineBackend._config_changed
is deprecated: use @observe and @unobserve instead.
```


```
    def _config_changed(self, name, old, new):
```

```
/Users/mkarroge/anaconda3/lib/python3.5/site-packages/traitlets/tra
itlets.py:770: DeprecationWarning: A parent of InlineBackend._confi
g_changed has adopted the new @observe(change) API
```

```
    clsname, change_or_name), DeprecationWarning)
```

Out[4]:

```
<matplotlib.collections.PolyCollection at 0x1057acc18>
```



Spectrogram

A spectrogram shows the spectrum of frequencies of sound over time. Here's a very pleasant video explaining how to read a spectrogram: <https://www.youtube.com/watch?v= FatxGN3vAM> (<https://www.youtube.com/watch?v= FatxGN3vAM>).


Here it is for this track:

In [5]:

```
X = librosa.stft(x)
Xdb = librosa.amplitude_to_db(abs(X))
plt.figure(figsize=(14, 5))
librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='hz')
#If to print log of frequencies
#librosa.display.specshow(Xdb, sr=sr, x_axis='time', y_axis='log')
plt.colorbar()
```

Out[5]:

<matplotlib.colorbar.Colorbar at 0x125848080>



That's a noisy song!

Zero-crossing rate

The rate of sign-changes along a signal; usually has higher values for highly percussive sounds (like in metal/rock)

In [6]:


```
x, sr = librosa.load(audio_path)
#Plot the signal:
plt.figure(figsize=(14, 5))
librosa.display.waveplot(x, sr=sr)
```

/Users/mkarroge/anaconda3/lib/python3.5/site-packages/librosa/core/audio.py:146: UserWarning: PySoundFile failed. Trying audioread instead.

```
warnings.warn('PySoundFile failed. Trying audioread instead.')
```


Out[6]:

<matplotlib.collections.PolyCollection at 0x125892e10>



In [7]:

```
# Zooming in
n0 = 9000
n1 = 9100
plt.figure(figsize=(14, 5))
plt.plot(x[n0:n1])
plt.grid()
```



We can see here that the signal crosses zero 5 times in the above graph, which looks at the first 100 columns in our 1D array. It can also be calculated as follows:

In [9]:

```
zero_crossings = librosa.zero_crossings(x[n0:n1], pad=False)
zero_crossings
```

Out[9]:

```
array([False, False, False, False, False, False,  True, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False,  True, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False])
```

In [10]:

```
# for entire song
zero_crossings = librosa.zero_crossings(x, pad=False)
print(sum(zero_crossings))
```

419596

In [13]:

```
zero_crossings = librosa.zero_crossings(x[n0:n1], pad=False)
print(sum(zero_crossings))
```

5

The zero crossing rate is represented in our dataset, and I think will be a good indicator of the "roughness" of a track (as high percussion is probably more likely to be rock)

Spectral Centroid

This feature shows where the "center of mass" for a sound is; it is the weighted mean of all present frequencies.

In [18]:

```
# spectral centroid
import sklearn
spectral_centroids = librosa.feature.spectral_centroid(x, sr=sr)[0]
spectral_centroids.shape # shape of array with column for each frame in sample
```

Out[18]:

```
(9402,)
```

In [21]:


```
# Computing the time variable for visualization
frames = range(len(spectral_centroids))
t = librosa.frames_to_time(frames)

# Normalizing viz
def normalize(x, axis=0):
    return sklearn.preprocessing.minmax_scale(x, axis=axis)

# Plotting the Spectral Centroid along the waveform
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_centroids), color='r')
```

Out[21]:

```
[<matplotlib.lines.Line2D at 0x12ff392e8>]
```




In [23]:

```
plt.plot(t, spectral_centroids, color='r') # without normalizing
```

Out[23]:

```
[<matplotlib.lines.Line2D at 0x1258f79b0>]
```



Mel-Frequency Cepstral Coefficient (MFCC)

These coefficients describe the "shape" of a sound signal. It can also be used to describe the timbre. Timbre, or tone is described with words like sharp, round, reedy, brassy, bright, etc.

In [26]:

```
mfccs = librosa.feature.mfcc(x, sr=sr)
print(mfccs.shape)
```

```
(20, 9402)
```


The tuple above gives us the total number of MFCCs calculated (index 0) and the total number of frames available.

In [32]:

```
#Displaying the MFCCs:  
librosa.display.specshow(mfccs, sr=sr, x_axis='time')
```

Out[32]:

<matplotlib.axes._subplots.AxesSubplot at 0x1248af3c8>




In [34]:

```
#Displaying the MFCCs:  
librosa.display.specshow(mfccs[1:20], sr=sr, x_axis='time')
```

Out[34]:

<matplotlib.axes._subplots.AxesSubplot at 0x127b9d4e0>



This one is without the first feature, which contains a constant offset; it is often discarded according to [this](https://musicinformationretrieval.com/mfcc.html) (<https://musicinformationretrieval.com/mfcc.html>).

In []: